



UNIVERSITÀ DI PISA

BLOOM FILTERS in MAP-REDUCE

Computer Engineering / Artificial Intelligence and Data Engineering

Cloud Computing

Fabiano Pilia

Riccardo Sagramoni

Emanuele Tinghi

Veronica Torraca

A.Y. 2021-2022

Table of Contents

1.	Introduction	4
1.1.	Project Specification	4
1.2.	Project Goal	5
2.	Pseudocode	6
2.1.	Mapper	6
2.2.	Reducer	6
2.3.	Tester	7
3.	Preparation of dataset and infrastructure	8
3.1.	Partitioning of the dataset	8
3.2.	Count the number of movies for each rating	8
3.3.	Infrastructure configuration	9
4.	Hadoop implementation	10
4.1.	IntArrayWritable and BooleanArrayWritable	10
4.2.	BloomFilterConfigurationName and MapReduceParameters	10
4.3.	BloomFilterUtil	11
4.4.	BloomFilterMapper	11
4.5.	BloomFilterReducer	14
4.6.	Tester	16
4.6.1.	Writables	17
4.6.2.	Implementation	17
5.	Spark implementation	21
5.1.	Bloom filters builder	21
5.1.1.	Input parameters	21
5.1.2.	Broadcast variables	21
5.1.3.	MapReduce job	21
5.2.	Bloom filters tester	24
5.2.1.	Input parameters	24
5.2.2.	Broadcast variables	24
5.2.3.	MapReduce job	24
5.3.	bloomfilters_util.py	27

5.3.1.	get_size_of_bloom_filters	27
5.3.2.	compute_indexes_to_set	27
5.3.3.	create_pair_rating_indexes	27
6.	Experimental results	29
6.1.	Hadoop	29
6.2.	Spark.....	30

1. Introduction

A bloom filter is a space efficient probabilistic structure used for membership testing. The aim of this project is to implement and design a bloom filter and the algorithm that keeps it updated both in Spark and in Hadoop.

1.1. Project Specification

Bloom filter, given their space efficient nature, are not perfect. Actually, it has a fixed size, so, independently of how many keys you will pass to it later, you even get to decide how much memory you want it to use. As already said, this comes at a certain cost, which is the probability of having false positives. But there can never be false negatives.




Figure 1: Bloom filter example.

A bloom filter is a bit-vector with m elements. It uses k hash functions to map n keys to the m elements of the bit-vector. Given a key id , every hash function (h_1, \dots, h_k) computes the output positions and sets the corresponding bit in that position to 1, if it is equal to 0.

Let's consider a bloom filter with the following characteristics:

- m : number of bits in the bit-vector
- k : number of hash functions
- n : number of keys added for membership testing
- p : false positive rate (probability between 0 and 1)

The relation between these values can be expressed as:

$$m = \frac{n * \ln (p)}{(\ln (2))^2}$$

$$k = \frac{m}{n} * \ln (2)$$

$$p \approx \left(1 - e^{-\frac{k*n}{m}}\right)^k$$

To design a bloom filter with a given false positive rate p , you need to estimate the number of keys n to be added to the bloom filter, then compute the number of bits m in the bloom filter and finally compute the number of hash functions k to use.

1.2. Project Goal

The goal of this project is to build a bloom filter over the rating of movies listed in this [IMDB dataset](#). The average ratings must be rounded to the closest integer value and every rating value must have its own bloom filter.

The bloom filter construction must be implemented following the MapReduce workflow, both in Hadoop and Spark frameworks. Moreover, the exact number of false positive for each rating must be computed.

In the Hadoop implementation the following classes must be followed:

- `org.apache.hadoop.mapreduce.lib.input.NLineInputFormat`: splits N lines of input as one split
- `org.apache.hadoop.util.hash.Hash.MURMUR_HASH`: the hash function family to use

In the Spark implementation analogous classes must be used/implemented.

2. Pseudocode

In our project we choose to design two types of MapReduce applications, both in Hadoop and Spark implementations:

- The **Builder** application: to build the actual Bloom filter for each rating
- The **Tester** application: to test the Bloom filter created in the Builder, computing the number of false positives

For this aim, we also randomly split the initial dataset in **Train** and **Test** partitions, using another Spark application we realized.

2.1. Mapper

The mapper receives a text line containing the tuple (movieId, rating, numberOfVotes), from which extracts the movieId and the rating, rounded to the closest integer. Then, for each movie computes the hashes of the ID, that correspond to the positions to set in the bloom filter and emits the key-value pair <rating, positions_list>.

```
3  class Mapper:
4      method SETUP()
5          BLOOM_FILTERS_SIZE <- {key_1: size_1, ..., key_m: size_m}
6          HASH_FUNCTION_NUMBER <- # of hash functions to compute
7      method MAP(line_id, line_text)
8          movieId <- first token from line_text
9          rating <- second token from line_text
10         for h = 0 to HASH_FUNCTION_NUMBER do
11             positions[i] <- hash(movieId, i) % BLOOM_FILTER_SIZE
12         EMIT(rating, positions[])
```

2.2. Reducer

The reducer receives several lists of positions to set to true in the bloom filter associated with the specific rating. It creates the bloom filter, sets to true the positions specified and emits the key-value pair <rating, bloom_filter>.

```
14  class Reducer:
15      method SETUP()
16          BLOOM_FILTERS_SIZE <- {bf_1: size_1, ..., bf_m: size_m}
17      method REDUCE(rating, list_of_positions)
18          for i = 0 to BLOOM_FILTERS_SIZE[rating] do
19              bloom_filter[i] <- false
20          for all positions in list_of_positions do
21              for all pos in positions do
22                  bloom_filter[pos] <- true
23          EMIT(rating, bloom_filter)
```

2.3. Tester

During the test phase, given a set of films that shouldn't be inside the bloom filter previously created, the Tester application computes the hashes of all the movieIds (as in the Builder Mapper) and then check, for each movie, if all the corresponding positions, i.e. the hash values, are set to true in the bloom filter. Every movie of the test dataset presents in the bloom filter is a false positive one and it is added to the total count.

The ReducerTester emits the key-value pair <rating, (number_of_false_positives, number_of_tests, false_positive_percentage)>.

```
27 class MapperTesterHash:
28     same as Mapper (builder)
```

```
33 class ReducerTester:
34     method REDUCE(rating, list_of_positions)
35         number_of_tests <- 0
36         number_of_false_positives <- 0
37         for all positions in list_of_positions do
38             if bloom_filter[pos] is true for all pos in positions
39                 number_of_false_positive <- number_of_false_positive + 1
40         number_of_tests <- number_of_tests + 1
41         false_positive_percentage <- number_of_false_positives / number_of_tests
42         EMIT(rating, (number_of_false_positives, number_of_tests, false_positive_percentage))
```

3. Preparation of dataset and infrastructure

3.1. Partitioning of the dataset

In order to test the bloom filters, the application **requires a set of records which weren't used during the building process**. So, we decided to split the dataset `imdb.tsv` into two partitions:

- `train_imdb.tsv`, used for the building application.
- `test_imdb.tsv`, used for the testing application.

The partitioning of the dataset has been implemented as a Spark application in the file `split-dataset.py`, in order to exploits both the benefits of MapReduce and the speed of Spark.

- Read the dataset from file
- Remove the first line, which contains the head of the TSV file
- Randomly split the lines into two RDDs
- Save the RDDs into text files

The final partition will be *approximately* (because of the random split) 60% for building and 40% for testing.

This script is executed only once when preparing the data for the project.

3.2. Count the number of movies for each rating

In order to estimate the size of the bloom filters, we require to know **how many movies exist for each rating value** (i.e. how many keys will be inserted in the bloom filter).

A Spark implementation of the process was provided (`count-number-of-keys.py`), in order to benefit from the task and data parallelization.

The application simply reads the “train” partition of the dataset, counts how many records exists for each rating (rounded to the closest integer) and outputs the results to a file in the HDFS.

The generated files (one for each executor) are merged into a single file by the `getmerge` HDFS command and reupload to HDFS, in order to be accessible by Java and Python when reading its content.

This script is executed only once when preparing the data for the project.

3.3. Infrastructure configuration

The infrastructure was configured with the settings for 2GB nodes, as presented during the course.

In order to fully exploit the cluster during the Spark execution, the number of executor instances has been set to 4 (equals to the available nodes in the cluster) in the `spark-defaults.conf` file.

The cluster we used consists of 4 nodes on as many machines, available at the ip addresses:

- 172.16.4.141 -> `hadoop-namenode`
- 172.16.4.163 -> `hadoop-datanode-2`
- 172.16.4.172 -> `hadoop-datanode-3`
- 172.16.4.210 -> `hadoop-datanode-4`

4. Hadoop implementation

For the Hadoop implementation of the Bloom filter, the classes and structures described in the following sections were realized.

Moreover, we chose to set the number of lines for the mapper (`lines_per_map` parameter) to **15000 lines** for each mapper, because, based on the dataset size, we obtained about 10 mappers for each node of the cluster, which is coherent with the Hadoop guidelines.

As regards the number of reducers, we chose to use **4 reducers** in order to exploit all the cluster nodes and we saw that empirically that is more efficient with respect the usage of 3 reducers, in fact we get the following results:

Launched reduce tasks=3

Total time spent by all reduces in occupied slots (ms)=293336

Total time spent by all reduce tasks (ms)=146668

Launched reduce tasks=4

Total time spent by all reduces in occupied slots (ms)=290606

Total time spent by all reduce tasks (ms)=145303

4.1. IntArrayWritable and BooleanArrayWritable

The classes `IntArrayWritable` and `BooleanArrayWritable` were created to properly handled the output values between the Mapper and the Reducer and between Builder and Tester applications.

These two new classes extend the `ArrayWritable` class and can be serialized and deserialized to allow Hadoop to use them inside the cluster.

4.2. BloomFilterConfigurationName and MapReduceParameters

These two structures were realized to store the configuration parameters both for the Bloom filters and the MapReduce implementation. In particular:

- **BloomFilterConfigurationName** is an enumerator type that provides the names of the properties for the configuration of the MapReduce application in order to avoid to use directly the parameter name inside the code
- **MapReduceParameters** is a Singleton class for reading the parameters from file for the MapReduce jobs configuration and contains information about the numbers of Reducers and lines received by a mapper (both for Builder and Tester application).

4.3. BloomFilterUtil

BloomFilterUtil is a class containing helpful methods used in the code, both in Builder and in Tester application, such as:

- **generateConfiguration**: creates the configuration for the MapReduce jobs, in particular, sets how many bloom filters must be created and the size each of them must have and how many hash functions must be used
- **readConfigurationBloomFilterSize**: reads the configuration of the bloom filters to create a map structure containing the key-value pair <rating, size> for each bloom filter
- **getBloomFiltersSizeParameters**: reads from file the number of input key for each bloom filter and compute the required size of the data structure, given as argument also the false positive rate to achieve

4.4. BloomFilterMapper

The **BloomFilterMapper** class receives as input the id of the line extracted from the text file (the key) and the text line itself (the value). After the parsing of the line, getting the movieId and the rating, it **computes the hash values** required and maps those values to the corresponding rating.

```
28 public class BloomFilterMapper extends Mapper<LongWritable, Text, ByteWritable, IntArrayWritable> {
29     // Logger
30     private static final Logger LOGGER = LogManager.getLogger(BloomFilterMapper.class);
31
32     // Number of hash functions that must be applied
33     private int HASH_FUNCTIONS_NUMBER;
34     // Size of each bloom filter
35     private Map<Byte, Integer> BLOOM_FILTER_SIZE;
36
37     // Array of IntWritable for the output value of the mapper
38     private final IntArrayWritable outputValue = new IntArrayWritable();
39     // ByteWritable for the output key
40     private final ByteWritable outputKey = new ByteWritable();
41
42     // Instance of the hash function MURMUR_HASH
43     private final Hash hash = Hash.getInstance(Hash.MURMUR_HASH);
```

Before the map function, a **setup** function is executed for loading from the context all the information needed to correctly apply the map method, such as the **BLOOM_FILTER_SIZE** and the **HASH_FUNCTIONS_NUMBER** parameters

```

47      @Override
48      public void setup (Context context) {
49          // Retrieve configuration
50          Configuration configuration = context.getConfiguration();
51
52          // Read size of the bloom filters
53          BLOOM_FILTER_SIZE = BloomFilterUtils.readConfigurationBloomFiltersSize(configuration);
54
55          // Read how many hash functions must be implemented
56          HASH_FUNCTIONS_NUMBER = context.getConfiguration().getInt(
57              BloomFilterConfigurationName.NUMBER_HASH.toString(),
58              defaultValue: -1
59          );
60          LOGGER.debug("Number of hash functions = " + HASH_FUNCTIONS_NUMBER);
61      }

```

As already mentioned, the map function emit the key-value pair <rating, list of hashes>, where each hash value corresponds to the position to set to true in the bloom filter, and the rating is the vote of a movie, rounded to the closest integer number.

```

65     @Override
66     public void map (LongWritable key, Text value, Context context)
67         throws IOException, InterruptedException
68     {
69         // Create string tokenizer to extract movie id and rating
70         StringTokenizer itr = new StringTokenizer(value.toString());
71
72         // Retrieve the tokens obtained
73         if (!itr.hasMoreTokens()) {
74             LOGGER.error("Input line has not enough tokens: " + value);
75             return;
76         }
77         String movieId = itr.nextToken();
78         LOGGER.debug("movieId = " + movieId);
79
80         if (!itr.hasMoreTokens()){
81             LOGGER.error("Input line has not enough tokens: " + value);
82             return;
83         }
84         byte rating = (byte) Math.round(Double.parseDouble(itr.nextToken()));
85         LOGGER.debug("Rating = " + rating);
86
87
88
89         Integer bloomFilterSize = BLOOM_FILTER_SIZE.get(rating);
90         if (bloomFilterSize == null) {
91             LOGGER.error("Rating key " + rating + " doesn't exist in linecount");
92             return;
93         }
94
95         // Local array of IntWritable to contain the hashes of the movie's id
96         IntWritable[] hashes = new IntWritable[HASH_FUNCTIONS_NUMBER];
97
98         // Apply the hash function with different seeds, to obtain HASH_FUNCTIONS_NUMBER values
99         for (int i = 0; i < HASH_FUNCTIONS_NUMBER; i++){
100             int hashValue = hash.hash(movieId.getBytes(StandardCharsets.UTF_8), i);
101
102             hashes[i] = new IntWritable(
103                 Math.abs(hashValue % bloomFilterSize)
104             );
105
106             LOGGER.debug("Computed hash n." + i + ": " + hashes[i]);
107         }
108
109         LOGGER.debug("WRITE (key, value) = ( " + rating + ", " + Arrays.toString(hashes) + " )");
110
111
112         // Setting the output values
113         outputKey.set(rating);
114         outputValue.set(hashes);
115
116
117         // Emit the key-value pair
118         context.write(outputKey, outputValue);
119     }

```

4.5. BloomFilterReducer

The **BloomFilterReducer** class receives the movie rating as key, and the list of positions to set to true in the Bloom filter.

```
24 public class BloomFilterReducer
25     extends Reducer<ByteWritable, IntArrayWritable, ByteWritable, BooleanArrayWritable>
26 {
27     // Logger
28     private static final Logger LOGGER = LogManager.getLogger(BloomFilterReducer.class);
29
30     // Size of each bloom filter
31     private Map<Byte, Integer> BLOOM_FILTER_SIZE;
32
33     // Writable array for the result of the reducer (i.e. the bloom filter)
34     private final BooleanArrayWritable SERIALIZABLE_BLOOM_FILTER = new BooleanArrayWritable();
```

In the setup method, **BLOOM_FILTER_SIZE** is loaded from the context, to correctly perform the reduce function.

```
38     @Override
39     public void setup (Context context) {
40         BLOOM_FILTER_SIZE = BloomFilterUtils.readConfigurationBloomFiltersSize(context.getConfiguration());
41         LOGGER.debug("Bloom filter size = " + BLOOM_FILTER_SIZE);
42     }
```

In the **reduce** method a temporary bloomFilter array of BooleanWritable is created and initialized with all the values to false. Then the arrays of positions received from the mapper are iterated and the corresponding indexes are set to true. After these steps, the reduce method returns the key-value pairs **<rating, bloomFilter>**, where the bloomFilter is the actual bloom filter array associated with the rating key.

```
46     @Override
47     public void reduce (ByteWritable key, Iterable<IntArrayWritable> values, Context context)
48         throws IOException, InterruptedException
49     {
50         LOGGER.debug("Reducer key = " + key);
51
52         // Instantiate the temporary bloom filter, i.e. an array of BooleanWritable
53         BooleanWritable[] bloomFilter = new BooleanWritable[BLOOM_FILTER_SIZE.get(key.get())];
54         for (int i = 0; i < bloomFilter.length; i++) {
55             bloomFilter[i] = new BooleanWritable(value: false);
56         }
```

```

60 // Iterate the intermediate data
61 for (IntArrayWritable array : values) {
62     // Generate an iterable array
63     IntWritable[] arrayWithHashedIndexes = (IntWritable[]) array.toArray();
64     LOGGER.debug("IntWritable array = " + Arrays.toString(arrayWithHashedIndexes));
65
66     // Iterate the list of BF indexes produced by mapper's hash functions
67     for (IntWritable i : arrayWithHashedIndexes) {
68         int indexToSet = i.get();
69         LOGGER.debug("indexToSet = " + indexToSet);
70
71         // Check if index is a valid number
72         if (indexToSet < 0 || indexToSet >= bloomFilter.length) {
73             LOGGER.warn("Index " + indexToSet +
74                 " for key " + key + " is not valid");
75             continue;
76         }
77
78         // Set to true the corresponding item of the array
79         if (!bloomFilter[indexToSet].get()) {
80             bloomFilter[indexToSet].set(true);
81         }
82     }
83 }
84
85 LOGGER.debug("Bloom filter length = " + bloomFilter.length);
86 LOGGER.debug("Bloom filter = " + Arrays.toString(bloomFilter));
87
88 // Emit the reducer's results
89 SERIALIZABLE_BLOOM_FILTER.set(bloomFilter);
90 context.write(key, SERIALIZABLE_BLOOM_FILTER);
91

```

4.6. Tester

In the **testing phase**, we have to **count how many false positives** there are in the generated bloom filters.

In order to achieve our goal, we have to do the following action during the map phase:

- *Parse* each record of the “test” partition of the dataset and compute the corresponding list of `indexes_to_set` in the bloom filter (this step is the same of the builder application, but it’s applied to a different input file)
- *Read* the generated bloom filters from file and distribute them to the reducers.

Since Hadoop lacks of an efficient method to distribute big amount of data between nodes (such as the *broadcasting variables* in Spark), we decided to distribute the bloom filters together with the arrays of `indexes_to_set` during the map phase.

Because of that we decided to implement **two different Mapper** classes (one for the array of `indexes_to_set`, while the other for the bloom filters) and to emit a **TesterGenericWritable** as map output value, i.e. a generic wrapper used to uniformly access all the values emitted by the mappers.

Moreover, we used the **MultipleInputs** class to support both files as mapper inputs (the text file containing the “test” dataset and the sequence file containing the bloom filters).

As regards the **reducer**, in order to properly compute the number of false positives, we needed to make sure that the bloom filter would be the first value in the **iterable structure** which contains the inputs. Thus, we implemented a **Secondary sorting** mechanism:

- We created a new key for the map output (**IntermediateKeyWritable**), which contains the rating of the movie and a boolean which indicates if the corresponding **TesterGenericWritable** is a Bloom filter (true) or an array of `indexes_to_set` (false)
- We modified the MapReduce **partitioner**, so that it would distribute the map output values only by the rating value
- We modified the MapReduce **key comparator**, so that Bloom filter structure would be the first in the list map output values for each map output key
- We modified the MapReduce **group comparator**, so that the map output values would be still grouped by the rating

4.6.1. Writables

These writable classes are used to wrap the values needed in the Tester application:

- **TesterGenericWritable**: extension of `GenericWritable`, is a wrapper for the output of the mappers
- **TesterResultsWritable**: implementation of `Writable`, contains the results of the Tester application as a tuple of *<false positive occurrences, total occurrences, false positive percentage>*
- **IntermediateKeyWritable**: implementation of a `WritableComparable`, is a wrapper for the intermediate values key between the two mappers and the reducer, to sort correctly the values before reaching the reducer.

```
13 public class TesterGenericWritable extends GenericWritable {
14
15     private static Class[] CLASSES = {
16         BooleanArrayWritable.class, // bloom filter
17         IntArrayWritable.class      // hash of the inputs (indexes to set)
18     };
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000
```

4.6.2. Implementation

The implementation for the Tester is like that of the MapReduce algorithm used in the Builder application of the Bloom filter. The major difference lays in the fact that there are two separate mapper classes:

- **MapperTesterForHashValues**: mapper used as in the builder mapper, to **compute the hashes** of the `movieId`, values corresponding to the positions to check if are set in the bloom filter, in order to count the exact number of false positive in the reduce phase. The difference from the Builder is the dataset file received in input, that contains movies not present in the other dataset, and the emitted value type, that in this case is yet an `IntArrayWritable`, but

wrapped in a `TesterGenericWritable` structure, for obtaining a **uniform output value type** from the second mapper tester

- **MapperTesterForBloomFilter**: mapper used to **distribute the bloom filters** generated in the builder to the right reducer task, in order to test if there are false positives. As in the first mapper tester, the output (the bloom filter) is wrapped in the `TesterGenericWritable` structure

About the **ReducerTester**, it receives in input both the Bloom filter previously generated and the list of `IntArrayWritable`, each of which represents the array containing the positions to check if set to true in the Bloom filter, for each movie in the test dataset.

To properly receive both mapper outputs, in order to have the Bloom filter as the first item in the input list of the reducer, and the arrays of positions after, we introduced a **composite key** `<rating, isBloomFilter>`, wrapped in the `IntermediateKeyWritable`, described earlier, and we implemented a Secondary Sort on that composite key through the classes

KeyComparator, GroupComparator, PartitionerTester:

```
16 class KeyComparator extends WritableComparator {
17
18     protected KeyComparator() { super(IntermediateKeyWritable.class, createInstances: true); }
19
20
21     @Override
22     public int compare(WritableComparable w1, WritableComparable w2) {
23         IntermediateKeyWritable keyWritable1 = (IntermediateKeyWritable) w1;
24         IntermediateKeyWritable keyWritable2 = (IntermediateKeyWritable) w2;
25
26         int cmp = Byte.compare(keyWritable1.getRating(), keyWritable2.getRating());
27         if (cmp != 0) {
28             return cmp;
29         }
30
31         return -Boolean.compare(keyWritable1.getIsBloomFilter(), keyWritable2.getIsBloomFilter()); // reverse order
32     }
33 }
34
35 }
```

```
10 class GroupComparator extends WritableComparator {
11     protected GroupComparator() { super(IntermediateKeyWritable.class, createInstances: true); }
12
13     @Override
14     public int compare(WritableComparable w1, WritableComparable w2) {
15         IntermediateKeyWritable keyWritable1 = (IntermediateKeyWritable) w1;
16         IntermediateKeyWritable keyWritable2 = (IntermediateKeyWritable) w2;
17         return Byte.compare(keyWritable1.getRating(), keyWritable2.getRating());
18     }
19 }
20 }
```

```
11 class PartitionerTester extends Partitioner<IntermediateKeyWritable, TesterGenericWritable> {
12     @Override
13     public int getPartition(IntermediateKeyWritable key,
14                             TesterGenericWritable value,
15                             int numReduceTasks)
16     {
17         return (key.getRating().hashCode() & Integer.MAX_VALUE) % numReduceTasks;
18     }
19 }
```

Then the actual **reduce** phase is very similar to that of the reducer used in the construction of the Bloom filters. The main difference is that the reducer in the Builder has to set the positions indicated to generate the Bloom filter, while the reducer Tester has to **check** if the positions, computed through the hash functions in the mapper for each test movie, are set or not in the Bloom filter. If at least one position is not set, then the movie is not at **false positive**, otherwise it is and has to be counted as such.

The output of this reducer is then composed by the key-value pair **<rating, (numberOfFalsePositives, totalTests, falsePositiveRate)>**, where the value of the pair is wrapped in a `TesterResultsWritable` structure.

```
35 class ReducerTester
36     extends Reducer<IntermediateKeyWritable, TesterGenericWritable, ByteWritable, TesterResultsWritable>
37 {
38     // Logger
39     private static final Logger LOGGER = LogManager.getLogger(ReducerTester.class);
40
41     private final ByteWritable outputKey = new ByteWritable();
42
43     // Result of tester execution
44     // (number of false positives, total number of test samples, false positive probability)
45     private final TesterResultsWritable testResults = new TesterResultsWritable();
46
47
48     @Override
49     public void reduce (IntermediateKeyWritable key, Iterable<TesterGenericWritable> values, Context context)
50         throws IOException, InterruptedException
51     {
52         LOGGER.debug("Reducer key = " + key);
53
54         // Get the bloom filter from the input values
55         TesterGenericWritable wrappedBloomFilter = values.iterator().next();
56         if (!(wrappedBloomFilter.get() instanceof BooleanArrayWritable)) {
57             LOGGER.error("BloomFilter " + key + " doesn't exist");
58             return;
59         }
60
61         BooleanWritable[] bloomFilter = (BooleanWritable[])
62             ( (BooleanArrayWritable) wrappedBloomFilter.get() )
63             .toArray();
64
65         LOGGER.debug("bloomFilter = " + Arrays.toString(bloomFilter));
66         LOGGER.debug("bloomFilter length = " + bloomFilter.length);
```

```

69     int numberOfTests = 0, numberOfFalsePositives = 0;
70
71     // Get the intermediate results from the mapper
72     for (TesterGenericWritable object : values) {
73
74         // Convert to array of IntWritable (the outputs of the hash functions)
75         IntWritable[] intArray = (IntWritable[]) ( (IntArrayWritable)object.get() ).toArray();
76         LOGGER.debug("intArray = " + Arrays.toString(intArray));
77
78         boolean isFalsePositive = true;
79
80         // Iterate the array of IntWritable in order to check the outputs of the hash functions
81         // (i.e. the position to hit in the bloom filter)
82         for (IntWritable i : intArray) {
83             int index = i.get();
84             LOGGER.debug("Index = " + index + " key = " + key +
85                 " BF_size = " + bloomFilter.length);
86
87             if (index < 0 || index >= bloomFilter.length) {
88                 LOGGER.error("Index " + index + " for key " + key +
89                     " not valid - out of bound");
90                 return;
91             }
92
93             // Check current value
94             if (!bloomFilter[index].get()) {
95                 // If AT LEAST ONE output is NOT set, then the sample is NOT a false positive
96                 isFalsePositive = false;
97                 break;
98             }
99         }
100     }
101
102     // Update statistics
103     numberOfTests++;
104     if (isFalsePositive) {
105         numberOfFalsePositives++;
106     }
107
108 }
109
110 // Set the results of the reducer
111 outputKey.set(key.getRating());
112 testResults.set(numberOfFalsePositives, numberOfTests);
113 context.write(outputKey, testResults);

```

5. Spark implementation

5.1. Bloom filters builder

The `bloomfilters_builder.py` file implements the building process of the Bloom filters in Spark. It reads the dataset and the `linecount` files from HDFS, compute the Bloom filter for each rating value and generate a *pickle file* containing them.

5.1.1. Input parameters

The following arguments are taken from command line:

- `false_positive_prob`: contains a float representing probability of false positive
- `dataset_input_file`: contains the path of the dataset used to build the Bloom filters as a string
- `linecount_file`: contains the path of the file containing the amount of the movies for each rating, used to compute the size of every single Bloom filter (see `bloomfilters_util.get_size_of_bloom_filters`)
- `output_file`: contains the path on which put the output file containing the bloom filters

5.1.2. Broadcast variables

We use three broadcast variables in order to efficiently distribute shared data between the worker nodes:

- `broadcast_hash_function_number`: contains the number of hashes to compute for each movie id.
- `broadcast_size_of_bloom_filters`: contains a dictionary that has as keys the ratings and as values the size of the corresponding bloom filter.

5.1.3. MapReduce job

To build the bloom filters using spark, the following steps are done:

- a. *Create* the broadcast variable `broadcast_hash_function_number`, containing the number of hashes to compute for each movie.
- b. *Create* the broadcast variable `broadcast_size_of_bloom_filters`, containing a dictionary that has as key the ratings and as values the size of the bloom filter of the corresponding rating.
- c. *Read and parallelize* the dataset using `textFile(dataset_input_file)`.
- d. *Parse* each line to extract the movie id and the average rating, round the average rating to the closest integer and compute a list with the hash values of the movie's id.
- e. **Group by rating** the lists computed previously and merge them into a unique list for each rating value, containing all the indexes of its bloom filter to be set to true (could contain duplicates).

- f. For each rating ***create a bloom filter*** by setting to true the indexes contained in the list obtained in the previous step.
- g. Save the result as a pickle file.

The following image shows the implementation of the code.

```
sc.textFile(dataset_input_file) \
    .map(lambda line:
        util.create_pair_rating_indexes(line,
                                         broadcast_size_of_bloom_filters.value,
                                         broadcast_hash_function_number.value)
    ) \
    .reduceByKey(extend_list) \
    .map(lambda pair_rating_indexes:
        (pair_rating_indexes[0],
         generate_bloom_filter(pair_rating_indexes[1],
                               broadcast_size_of_bloom_filters.value.get(pair_rating_indexes[0])))
    ) \
    .saveAsPickleFile(output_file)
```

5.1.3.1. *extend_list*

This function requires as input two lists and gives as output the first list extended with the values of the second list.

```
def extend_list(list1: list, list2: list) -> list:
    """
        Concat the second list to the first one and return the first one

        :param list1: first list
        :param list2: second list

        :return: merged list
    """
    list1.extend(list2)
    return list1
```

5.1.3.2. *generate_bloom_filter*

Given as input a list of `indexes_to_set` for a particular rating and the size of the Bloom filter, this function generates an array of that given size, at first with all the elements initialized to False

and then, setting to True each position of the array contained into the `indexes_to_set` list. The returned value is such array that represents the Bloom filter

The implementation can be seen in the next image.

```
def generate_bloom_filter(indexes_to_set: list, size: int) -> list:
    """
    Set to True the corresponding item of the Bloom Filter

    :param indexes_to_set: list of indexes to set True (computed through hash functions)
    :param size: size of the bloom filter

    :return: the Bloom filters structure
    """
    if size is None:
        return []

    bloom_filter = [False for _ in range(size)]

    for i in indexes_to_set:
        bloom_filter[i] = True

    return bloom_filter
```

5.2. Bloom filters tester

The **Bloom filter tester** application reads the *generated bloom filters* and the *test partition* of the dataset from HDFS and computes:

- The **number of false positives**
- The **total number of test records**
- The **false positive percentage**, defined as the ratio between the previous two values

5.2.1. Input parameters

The application requires the following console parameters:

- `false_positive_prob`: a float containing the desired false positive probability.
- `dataset_input_file`: the path to the test partition of the dataset.
- `bloom_filters_input_file`: the path to the file which contains the bloom filters.
- `linecount_file`: the path to the file containing the number of movies for each rating value, in the original training dataset.
- `output_file`: the file in which the output records must be stored.

5.2.2. Broadcast variables

We use three broadcast variables in order to efficiently distribute shared data between the worker nodes:

- `broadcast_hash_function_number`: contains the number of hashes to compute for each movie id.
- `broadcast_size_of_bloom_filters`: contains a dictionary that has as keys the ratings and as values the size of the corresponding bloom filter.
- `broadcast_bloom_filters`: contains all the bloom filters, generated by the **builder** application, as a dictionary object `{key: rating, value: bloom_filter}`.

5.2.3. MapReduce job

The tester is implemented as a MapReduce job with one parallelization, four transformations and one action.

1. *Read* the test partition of the dataset and parallelize it
2. A first **map transformation** generates a k-v pair consisting of the rating value and the list of indexes to set to true in the bloom filter


```
sc.textFile(dataset_input_file) \
    .map(lambda line:
        util.create_pair_rating_indexes(line,
                                         broadcast_size_of_bloom_filters.value,
                                         broadcast_hash_function_number.value)
    ) \
```

3. A second **map transformation** emits a k-v pair containing as key the rating value and as value a list containing:

- i. **1** or **0** if the testing record is a false positive or not (checked via the function `check_false_positive`, described later)
- ii. **1**, used to compute the number of occurrences for the current rating

```
.map(lambda rating_indexes: (rating_indexes[0],
                             [check_false_positive(rating_indexes[1],
                                                    broadcast_bloom_filters.value[rating_indexes[0]]),
                              1]
                             )
    ) \
```

4. The k-v pairs obtained in the previous map transformation are **reducedByKey**. The grouped values are the lists containing [1 or 0, 1]. These lists are summed element-wise, in order to parallelly compute the number of false positives and the total number of test for each rating

```
.reduceByKey(sum_elem_lists) \

def sum_elem_lists(list1: list, list2: list) -> list:
    list1 = [x + y for x, y in zip(list1, list2)]
    return list1
```

5. The last **map** stage emits a k-v pair containing as key the rating and as value a tuple composed by, in order, the number of false positives, the number of tested movies (with that rating) and the false positive rate

```
.map(lambda rating_counters: (rating_counters[0], (rating_counters[1][0],
                                                    rating_counters[1][1],
                                                    rating_counters[1][0] / rating_counters[1][1])
    )
    ) \
```

6. The result is saved using `saveAsTextFile` action

5.2.3.1. *check_false_positive*

To count the false positive occurrences, the function `check_false_positive` is used. This function requires as input the indexes of the Bloom filter to be checked and the Bloom filter itself.

The function iterates the list of indexes. If at least one index is equals to false, it means that the movie is not a false positive and return 0. Otherwise, if all the indexes are equals to true, the function returns 1.

```
def check_false_positive(indexes: list, bloom_filter: list) -> int:

    for i in indexes:
        if not bloom_filter[i]:
            return 0

    return 1
```

5.3. bloomfilters_util.py

This module defines some **utility functions** used by the two Spark applications. We will briefly discuss only the interesting ones.

5.3.1. get_size_of_bloom_filters

Generate a MapReduce job that reads the file with the number of keys for each rating in the “training” dataset and **compute the size of bloom filter for each rating**.

```
def get_size_of_bloom_filters(sc: SparkContext, linecount_file: str, false_positive_prob: float) -> dict:
    size_of_bf_list = sc.textFile(linecount_file) \
        .map(lambda line: line.split('\t')[0:2]) \
        .map(lambda split_line: (int(split_line[0]),
                                compute_size_of_bloom_filter(false_positive_prob, int(split_line[1]))))
        .collect()

    # Convert the list of tuples into a dictionary
    return dict(size_of_bf_list)
```

5.3.2. compute_indexes_to_set

It generates a list containing the indexes to set in the bloom filter, by computing the hash values of the movieid and by applying the modulo operation on them (so that it will be in the range of the bloom_filter_size).

The hash values are obtained by applying the MurmurHash algorithm with different seeds (the seed is equal to the position).

```
def compute_indexes_to_set(movie_id: str, rating: int, size_of_bloom_filters: dict, hash_function_number: int) -> list:

    # Get bloom filter size
    bloom_filter_size = size_of_bloom_filters.get(rating)

    if bloom_filter_size is None:
        return []

    # Compute hash values
    return [mmh3.hash(movie_id, i) % bloom_filter_size for i in range(hash_function_number)]
```

5.3.3. create_pair_rating_indexes

Given a line from the dataset, it parses the line to obtain the rating and the movie’s id and generate a tuple with the rating and a list with the indexes to set in the bloom filter (computed by calling the function compute_indexes_to_set)

```
def create_pair_rating_indexes(line: str, size_of_bloom_filters: dict, hash_function_number: int) -> Tuple[int, list]:

    # Split line and parse arguments
    split_line = line.split('\t')[0:2]
    movie_id = split_line[0]
    rating = int(float(split_line[1]) + 0.5) # round() is bugged in Python 3.6!!!

    return (
        rating,
        compute_indexes_to_set(movie_id,
                                rating,
                                size_of_bloom_filters,
                                hash_function_number)
    )
```

6. Experimental results

In order to standardize and automatize the execution of the application, **Linux shell scripts** were provided.

6.1. Hadoop

The following images contains, respectively, the output of the hadoop tester and the result of three different tests using different values of false positive probability.

```
4      (1825, 17337, 0.1052661936897964)
8      (14357, 140733, 0.10201587403096644)
2022-07-06 13:43:07,388 INFO sasl.SaslDataTra
check: localhostTrusted = false, remoteHostTr
1      (104, 972, 0.10699588477366255)
5      (4229, 40836, 0.10356058379860907)
9      (4630, 45182, 0.10247443672258864)
2022-07-06 13:43:07,393 INFO sasl.SaslDataTra
check: localhostTrusted = false, remoteHostTr
2      (266, 2581, 0.1030608291359938)
6      (8945, 87646, 0.10205827989868334)
10     (693, 6463, 0.10722574655732632)
3      (796, 7226, 0.1101577636313313)
7      (15118, 148449, 0.10183968905145875)

BLOOM FILTER TESTER COMPLETED
```

Rating	False Positive Probability		
	0.01	0.05	0.1
1	0.00925926	0.04629629	0.10699588
2	0.00891127	0.05695466	0.10306082
3	0.00899529	0.04926653	0.11015776
4	0.00911345	0.05012401	0.10526619
5	0.01021158	0.05135174	0.10356058
6	0.01014307	0.04979120	0.10205827
7	0.00987544	0.05103436	0.10183968
8	0.00983422	0.05167942	0.10201587
9	0.01002612	0.05307423	0.10247443
10	0.01299706	0.04487080	0.10722574
avg	0,00993668	0,05044432	0,10446552

6.2. Spark

The following pictures contain the output of the spark tester and the results of three tests.

```
(6, (8886, 87646, 0.10138511740410287))
(4, (1808, 17337, 0.10428563188556267))
(8, (14400, 140733, 0.10232141715162756))
(2, (250, 2581, 0.09686168151879117))
(10, (662, 6463, 0.10242921244004333))
2022-07-06 11:08:41,138 INFO sasl.SaslData
  false, remoteHostTrusted = false
(9, (4698, 45182, 0.10397946084724005))
(5, (4251, 40836, 0.10409932412577137))
(3, (792, 7226, 0.10960420703016883))
(7, (15263, 148449, 0.10281645548302784))
(1, (97, 972, 0.09979423868312758))

BLOOM FILTER TESTER COMPLETED
```

Rating	False Positive Probability		
	0.01	0.05	0.1
1	0.01028806	0.05452674	0.09979423
2	0.00929872	0.04533126	0.09686168
3	0.01093274	0.05175754	0.10960420
4	0.00911345	0.05075849	0.10428563
5	0.00999118	0.05066607	0.10409932
6	0.01039408	0.04972274	0.10138511
7	0.01024594	0.05079859	0.10281645
8	0.01002607	0.05057804	0.10232141
9	0.00960559	0.04964366	0.10397946
10	0.01005724	0.04858424	0.10242921
avg	0.00999531	0.05023674	0.10275767